

 [learn-co-curriculum](#) / [dsc-corr-autocorr-in-time-series-lab](#) Public [View license](#) 1 star  151 forks Star Watch ▾[Code](#) [Issues](#) [Pull requests](#) [Actions](#) [Projects](#) [Security](#) [Insights](#) solution ▾

...

This branch is [4 commits ahead](#), [5 commits behind](#) master.

Cheffrey2000 fixed env conflicts ...

on Oct 21, 2020

 8[View code](#) README.md

# Correlation and Autocorrelation in Time Series - Lab

## Introduction

In this lab, you'll practice your knowledge of correlation, autocorrelation, and partial autocorrelation by working on three different datasets.

## Objectives

In this lab you will:

- Plot and discuss the autocorrelation function (ACF) for a time series
- Plot and discuss the partial autocorrelation function (PACF) for a time series

# The Exchange Rate Data

---

We'll be looking at the exchange rates dataset again.

- First, run the following cell to import all the libraries and the functions required for this lab
- Then import the data in 'exch\_rates.csv'
- Change the data type of the 'Frequency' column
- Set the 'Frequency' column as the index of the DataFrame

```
# Import all packages and functions
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
from statsmodels.graphics.tsaplots import plot_acf, plot_pacf
from matplotlib.pyplot import rcParams
```

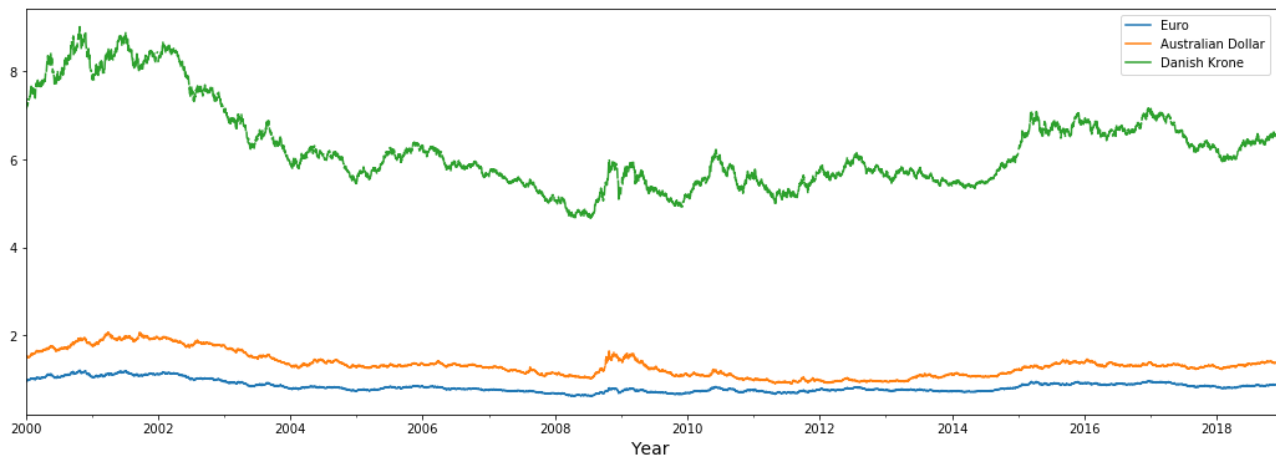
```
# Import data
xr = pd.read_csv('exch_rates.csv')

# Change the data type of the 'Frequency' column
xr['Frequency'] = pd.to_datetime(xr['Frequency'])

# Set the 'Frequency' column as the index
xr.set_index('Frequency', inplace=True)
```

Plot all three exchange rates in one graph:

```
# Plot here
xr.plot(figsize=(18,6))
plt.xlabel('Year', fontsize=14);
```



You can see that the EUR/USD and AUD/USD exchange rates are somewhere between 0.5 and 2, whereas the Danish Krone is somewhere between 4.5 and 9. Now let's look at the correlations between these time series.

```
# Correlation
xr.corr()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```

```
</style>
```

	Euro	Australian Dollar	Danish Krone
Euro	1.000000	0.883181	0.999952
Australian Dollar	0.883181	1.000000	0.882513
Danish Krone	0.999952	0.882513	1.000000

**What is your conclusion here? You might want to use outside resources to understand what's going on.**


```
# The exchange rates for Euro and the Australian dollar are highly correlated,
# but there are differences. The Euro and the Danish Krone, however, is perfectly cc
```

```
# If you do further research you'll notice that the Danish Krone is pegged to the Eu
# which means that they are basically designed to perfectly correlate together!
# The fact that the value is just very, very close to 1 is due to rounding errors.
# Usually when the correlation is so close to 1 (or -1), it's too good to be true.
# So make sure you always dig deeper to correctly understand and interpret these num
```

Next, look at the plots of the differenced (1-lag) series. Use subplots to plot them rather than creating just one plot.

```
# 1-lag differenced series
xr_diff = xr.diff(periods=1)

# Plot
xr_diff.plot(figsize=(13,8), subplots=True, legend=True);
```



Calculate the correlation of this differenced time series.

```
# Correlation
xr_diff.corr()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```

</style>

	Euro	Australian Dollar	Danish Krone
Euro	1.000000	0.545369	0.999667
Australian Dollar	0.545369	1.000000	0.545133
Danish Krone	0.999667	0.545133	1.000000

## Explain what's going on

```
# Differencing the series here led to a decrease
# in correlation between the EUR/USD and AUD/USD series.
# If you think a little further, this makes sense: in the previous lesson,
# the high correlation was a result of seasonality.
# Differencing led to an increase in correlation between series,
# here the series are moving in (more or less) the same direction
# on a day-to-day basis and seasonality is not present, hence this result.
```

Next, let's look at the "lag-1 autocorrelation" for the EUR/USD exchange rate.

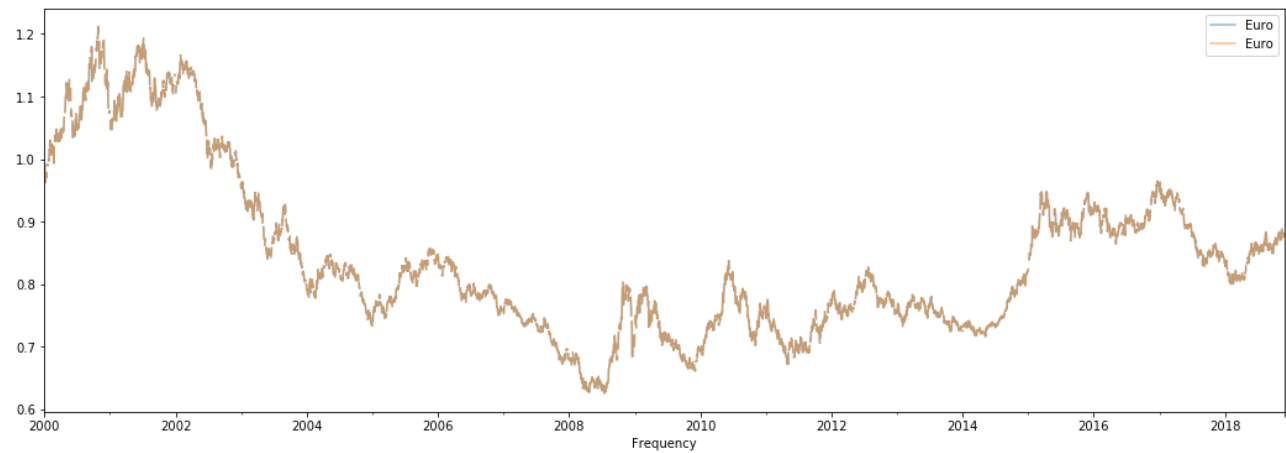
- Create a "lag-1 autocorrelation" series
- Combine both the original and the shifted ("lag-1 autocorrelation") series into a DataFrame
- Plot these time series, and look at the correlation coefficient

```
# Isolate the EUR/USD exchange rate
eur = xr[['Euro']]
```

```
# "Shift" the time series by one period
eur_shift_1 = eur.shift(periods=1)
```

```
# Combine the original and shifted time series
lag_1 = pd.concat([eur_shift_1, eur], axis=1)
```

```
# Plot
lag_1.plot(figsize=(18,6), alpha=0.5);
```



```
# Correlation
lag_1.corr()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }

.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}

</style>
```


	Euro	Euro
Euro	1.000000	0.999146
Euro	0.999146	1.000000

Repeat this for a "lag-50 autocorrelation".

```
# "Shift" the time series by 50 periods
eur_shift_50 = eur.shift(periods=50)

# Combine the original and shifted time series
lag_50 = pd.concat([eur_shift_50, eur], axis=1)
```

```
# Plot
lag_50.plot(figsize=(18,6), alpha=0.8);
```



```
# Correlation
lag_50.corr()
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```

```
</style>
```

	Euro	Euro
Euro	1.000000	0.968321
Euro	0.968321	1.000000


What's your conclusion here?

```
# Autocorrelation is very high in these time series, even up to a lag as big as 50!
# This is no big surprise though: remember that these are random walk series,
# which are highly recursive, as each value depends heavily on the previous one!
```



Knowing this, let's plot the ACF now.


```
# Plot ACF  
plt.figure(figsize=(12,5))  
pd.plotting.autocorrelation_plot(eur.dropna());
```



The series is heavily autocorrelated at first, and then there is a decay. This is a typical result for a series that is a random walk, generally you'll see heavy autocorrelations first, slowly tailing off until there is no autocorrelation anymore.

Next, let's look at the partial autocorrelation function plot.

```
# Plot PACF  
rcParams['figure.figsize'] = 14, 5  
plot_pacf(eur.dropna(), lags=50);
```





This is interesting! Remember that *Partial Autocorrelation Function* gives the partial correlation of a time series with its own lagged values, controlling for the values of the time series at all shorter lags. When controlling for 1 period, the PACF is only very high for one-period lags, and basically 0 for shorter lags. This is again a typical result for random walk series!

## The Airpassenger Data

Let's work with the air passenger dataset you have seen before. Plot the ACF and PACF for both the differenced and regular series.

Note: When plotting the PACF, make sure you specify `method='ywm'` in order to avoid any warnings.

```
# Import and process the air passenger data
air = pd.read_csv('passengers.csv')
air['Month'] = pd.to_datetime(air['Month'])
air.set_index('Month', inplace=True)
air.head()
```

<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }


```
.dataframe tbody tr th {
    vertical-align: top;
}

.dataframe thead th {
    text-align: right;
}
```


</style>

	#Passengers
Month	
1949-01-01	112
1949-02-01	118
1949-03-01	132
1949-04-01	129
1949-05-01	121

```
# Plot ACF (regular)
plt.figure(figsize=(12,5))
pd.plotting.autocorrelation_plot(air);
```




```
# Plot PACF (regular)
rcParams['figure.figsize'] = 14, 5
plot_pacf(air.dropna(), lags=50, method='ywm');
```




```
# Generate a differenced series
air_diff = air.diff(periods=1)
```

```
# Plot ACF (differenced)
plt.figure(figsize=(12,5))
pd.plotting.autocorrelation_plot(air_diff.dropna());
```



```
# Plot PACF (differenced)
rcParams['figure.figsize'] = 14, 5
plot_pacf(air_diff.dropna(), lags=50, method='ywm');
```



## Your conclusion here

```
# The result reminds us a lot of the google trends data.
# The seasonality is much more clear in the differenced time series.
# The PACF has just one very strong correlation, right at 12 months.
```

## The NYSE data

Are you getting the hang of interpreting ACF and PACF plots? For one final time, plot the ACF and PACF for both the NYSE time series.

Note: When plotting the PACF, make sure you specify `method='ywm'` in order to avoid any warnings.

```
# Import and process the NYSE data
nyse = pd.read_csv('NYSE_monthly.csv')
nyse['Month'] = pd.to_datetime(nyse['Month'])
nyse.set_index('Month', inplace=True)
nyse.head()
```

<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }


```
.dataframe tbody tr th {
    vertical-align: top;
}
```

```
.dataframe thead th {
    text-align: right;
}
```


</style>

	monthly_return
Month	
1961-01-01	0.082
1961-02-01	0.059
1961-03-01	0.050
1961-04-01	0.008
1961-05-01	0.042

```
plt.figure(figsize=(12,5))
pd.plotting.autocorrelation_plot(nyse.dropna());
```




```
rcParams['figure.figsize'] = 14, 5
plot_pacf(nyse, lags=30, method='ywm');
```



## Your conclusion here

```
# Autocorrelations and partial autocorrelations are virtually 0 for any lag.
# This is no surprise! The NYSE series was a white noise series, meaning there is no
# This is, again, a typical result for these kind of series.
```



## Summary

Great, you've now been introduced to ACF and PACF. Let's move into more serious modeling with autoregressive and moving average models!

## Releases

No releases published

---

## Packages

No packages published

---

## Contributors 4



**LoreDirick** Lore Dirick



**sumedh10** Sumedh Panchadhar



**mas16** matt



**Cheffrey2000** Jeffrey Hinkle

---

## Languages

● **Jupyter Notebook** 100.0%