

ARMA Models in StatsModels

Introduction

In this lesson, you'll use your knowledge of the autoregressive (AR) and moving average (MA) models, along with the `statsmodels` library to model time series data.

Objectives

You will be able to:

- Fit an AR model using `statsmodels`
- Fit an MA model using `statsmodels`

Generate a first order AR model

Recall that the AR model has the following formula:

$$Y_t = \mu + \phi * Y_{t-1} + \epsilon_t$$

This means that:

$$Y_1 = \mu + \phi * Y_0 + \epsilon_1$$

$$Y_2 = \mu + \phi * (\text{mean-centered version of } Y_1) + \epsilon_2$$

and so on.

Let's assume a mean-zero white noise with a standard deviation of 2. We'll also create a daily datetime index ranging from January 2017 until the end of March 2018.

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import numpy as np

np.random.seed(11225)

# Create a series with the specified dates
dates = pd.date_range('2017-01-01', '2018-03-31')
```


We will generate a first order AR model with $\phi = 0.7$, $\mu = 5$, and $Y_0 = 8$.

```
In [2]: error = np.random.normal(0, 2, len(dates))
Y_0 = 8
mu = 5
phi = 0.7
```

```
In [3]: TS = [None] * len(dates)
y = Y_0
for i, row in enumerate(dates):
    TS[i] = mu + y * phi + error[i]
    y = TS[i] - mu
```

Let's plot the time series to verify:

```
In [4]: series = pd.Series(TS, index=dates)
series.plot(figsize=(14,6), linewidth=2, fontsize=14);
```



Look at the ACF and PACF of the model


Although we can use `pandas` to plot the ACF, we highly recommended that you use the `statsmodels` variant instead.

```
In [5]: from statsmodels.graphics.tsaplots import plot_pacf
from statsmodels.graphics.tsaplots import plot_acf

fig, ax = plt.subplots(figsize=(16,3))
plot_acf(series, ax=ax, lags=40);

fig, ax = plt.subplots(figsize=(16,3))
plot_pacf(series, ax=ax, lags=40);
```

/opt/saturncloud/envs/saturn/lib/python3.10/site-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.
warnings.warn()



Check the model with ARMA in statsmodels

statsmodels also has a tool that fits ARMA models to time series. The only thing we have to do is provide the number of orders for AR and MA. Have a look at the code below, and the output of the code.

The `ARIMA()` function requires two arguments:

1. The first is the time series to which the model is fit
2. The second is the `order` of the model in the form (p, d, q)
 - `p` refers to the order of AR
 - `d` refers to the order of I (which will be discussed in a future lesson -- for now just use 0)
 - `q` refers to the order of MA

For example, a first order AR model would be represented as $(1, 0, 0)$.

```
In [6]: # Import ARIMA
from statsmodels.tsa.arima.model import ARIMA
import statsmodels.api as sm

# Instantiate an AR(1) model to the simulated data
mod_arma = ARIMA(series, order=(1,0,0))
```

Once you have instantiated the AR(1) model, you can call the `.fit()` method to fit the model to the data.

```
In [7]: # Fit the model to data
res_arma = mod_arma.fit()
```

Similar to other models, you can then call the `.summary()` method to print the information of the model.

```
In [8]: # Print out summary information on the fit
print(res_arma.summary())
```

```
SARIMAX Results
=====
Dep. Variable:          y      No. Observations:          455
Model:                ARIMA(1, 0, 0)      Log Likelihood          -968.698
Date:                Sat, 19 Nov 2022      AIC                  1943.395
Time:                12:13:39      BIC                  1955.756
Sample:              01-01-2017      HQIC                 1948.265
                  - 03-31-2018
Covariance Type:      opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
const          4.9664          0.262      18.937      0.000          4.452          5.480
ar.L1           0.6474          0.036      18.108      0.000          0.577          0.718
sigma2          4.1327          0.289      14.289      0.000          3.566          4.700
=====
Ljung-Box (L1) (Q):                0.00      Jarque-Bera (JB):                0.99
Prob(Q):                          0.99      Prob(JB):                0.61
Heteroskedasticity (H):            1.02      Skew:                    0.08
Prob(H) (two-sided):              0.89      Kurtosis:                2.84
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).

Make sure that the output for the ϕ parameter and μ is as you'd expect. You can use the `.params` attribute to check these values.

```
In [9]: # Print out the estimate for the constant and for theta
print(res_arma.params)
```

```
const      4.966442
ar.L1       0.647429
sigma2      4.132665
dtype: float64
```

Generate a first order MA model

Recall that the MA model has the following formula:

$$Y_t = \mu + \epsilon_t + \theta * \epsilon_{t-1}$$

This means that:

$$Y_1 = \mu + \epsilon_1 + \theta * \epsilon_0$$

$$Y_2 = \mu + \epsilon_2 + \theta * \epsilon_1$$

and so on.

Assume a mean-zero white noise with a standard deviation of 4. We'll also generate a daily datetime index ranging from April 2015 until the end of August 2015.

We will generate a first order MA model with $\theta = 0.9$ and $\mu = 7$.

```
In [10]: np.random.seed(1234)


# Create a series with the specified dates
dates = pd.date_range('2015-04-01', '2015-08-31')

error = np.random.normal(0, 4, len(dates))
mu = 7
theta = 0.9

TS = [None] * len(dates)
error_prev = error[0]
for i, row in enumerate(dates):
    TS[i] = mu + theta * error_prev + error[i]
    error_prev = error[i]
```

```
In [11]: series = pd.Series(TS, index=dates)

series.plot(figsize=(14,6), linewidth=2, fontsize=14);
```




Look at the ACF and PACF of the model

```
In [12]: fig, ax = plt.subplots(figsize=(16,3))
plot_acf(series, ax=ax, lags=40);

fig, ax = plt.subplots(figsize=(16,3))
plot_pacf(series, ax=ax, lags=40);
```

/opt/saturncloud/envs/saturn/lib/python3.10/site-packages/statsmodels/graphics/tsaplots.py:348: FutureWarning: The default method of 'yw' can produce PACF values outside of the [-1,1] interval. After 0.13, the default will change to unadjusted Yule-Walker ('ywm'). You can use this method now by setting method='ywm'.
warnings.warn(



Check the model with ARMA in statsmodels

Let's fit an MA model to verify the parameters are estimated correctly. The first order MA model would be represented as $(0, 1)$.

```
In [13]: # Instantiate and fit an MA(1) model to the simulated data
mod_arma = ARIMA(series, order=(0,0,1))
res_arma = mod_arma.fit()

# Print out summary information on the fit
print(res_arma.summary())
```

```
SARIMAX Results
=====
Dep. Variable:          y      No. Observations:         153
Model:                ARIMA(0, 0, 1)    Log Likelihood        -426.378
Date:                 Sat, 19 Nov 2022    AIC                   858.757
Time:                 12:22:11           BIC                   867.848
Sample:               04-01-2015         HQIC                  862.450
                   - 08-31-2015
Covariance Type:      opg
=====
              coef    std err          z      P>|z|      [0.025    0.975]
-----
const          7.5373      0.625     12.069      0.000        6.313      8.761
ma.L1          0.8727      0.047     18.489      0.000        0.780      0.965
sigma2        15.2765      1.592      9.597      0.000       12.156     18.397
=====
Ljung-Box (L1) (Q):                5.72    Jarque-Bera (JB):                9.11
Prob(Q):                           0.02    Prob(JB):                  0.01
Heteroskedasticity (H):              1.14    Skew:                      -0.48
Prob(H) (two-sided):                 0.65    Kurtosis:                   3.70
=====
```

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```
In [14]: # Print out the estimate for the constant and for theta  
print(res_arma.params)
```

```
const      7.537262  
ma.L1      0.872686  
sigma2     15.276484  
dtype: float64
```

Summary

Great job! In this lesson, you saw how you can use the AR and MA models using the `ARIMA()` function from `statsmodels` by specifying the order in the form of (p, q) , where at least one of p or q was zero depending on the kind of model fit. You can use `ARIMA()` to fit a combined ARMA model as well -- which you will do in the next lab!